

Javascript Language Extensions

Esha Choukse, Lara Schmidt

I. INTRODUCTION AND PROJECT GOALS

V8 is Google's compiler for javascript that is designed to make javascript as fast as possible. It has two modes of operation: Full compiler and Crankshaft compiler (Optimized). It compiles the code at the granularity of functions. It has many heuristics to decide between when to shift from unoptimized to optimized compilation. However one thing that V8 does not provide is, a way for a developer or user to provide 'hints' to the compiler. Since it dynamically compiles the code, the compiler doesn't have a full picture of the code. Hence, programmer is always in a better position to guide the compiler, as to when optimization is really needed. Therefore, by allowing the compiler to take hints about whether or not it should optimize a function, we can allow V8 to gain even more speed.

II. API

We implemented three optimizations to prototype our ideas. A developer could specify the optimization to use by adding code into the javascript. The developer defines an optimization per function. Again, this is because V8's optimization and code-tracking is done per function.

The first optimization is ***Optimize Immediately***. This optimization will optimize a function immediately upon it's creation. If a user knows that a function is going to be used a lot of times, why wait for the compiler to realize it. The second optimization is ***Never Optimize***. This will cause a function to never attempt to be optimized. This can be used if the function has volatile operations that we don't want to be re-ordered. This is also useful if we know that the function just satisfies some optimization condition, but is never run after being optimized. The third optimization is ***Edge Weight Profiling***. This one is more complicated and has two parts (and two corresponding user inputs). We will shortly explain the details on this optimization.

The programmer can specify these optimizations inside a multi-line comment with a special header. They specify the function name and the respective optimization. An int was used to make parsing easier. 0 for Optimize Immediately, 1 for Never optimize and 2 and 3 for Edge Weight Profiling. The two flags for Edge Weight Profiling have different user inputs associated with them. For example:

```
/*LEZ  
fname1 , 0 ;  
fname2 , 2 , 100 ;  
fname2 , 3 , 5 ;
```

```
fname3 , 1 ;  
*/
```

In this example fname will be optimized immediately, fname2 will have the Edge Weight Profiling optimization with input of 100 and 5, and fname3 will be never optimized. One could give several optimizations to one function by just putting the function name in several lines as with fname2.

We also wanted to provide an ability for the user to optimize the javascript without the developer console. So, we created a Chrome browser extension that would insert the multi-line code into a viewed html file and cause our optimizations to be run. The browser inserts the javascript comments into a script tag at the beginning of html pages before the page is run, allowing it to be reached first by V8. For a prototype we just allowed the user to have a set of saved optimizations that was then inserted into pages on a certain domain, however one can see how it could be extended to allow the user to specify different optimizations for different pages.

III. EDGE WEIGHT PROFILING

As promised, let us dig deeper into the Edge Weight Profiling optimization. While, ideally we would like to have a better user-facing optimization, we ran into trouble with the complexity of V8. However this does give the developer more power. By default Chrome will mark a function for self optimization after certain conditions are hit. For example it runs X times or there is a loop with Y runs in it. The way it does this is to set a counter to a set value. By default it is set to 6144 and this, as well as all other constants we mention have a modifiable whole-system V8 flag. The counter decrements whenever a function returns (to count the times a function is run) and when a back-edge of a loop is taken.

The weight of the default return decrement is set to 6144/130. This will for example cause a function to optimize after 130 iterations. Our first part of this optimization allows the user to set the 130 to something else. Therefore if a user sets the number lower, it will optimize quicker.

The loop counter by default decrements by 1 or by *amount of code in loop / 170*, whichever is larger. Our second optimization allows the user to modify the weight of this counter. Therefore a larger weight will optimize faster.

However though this logic is very evident in the code, we sometimes see some weird behaviour from Chrome that we have not been able to figure out within limited time. For example Chrome optimizes a function that

runs 1000 times. But if it has a while loop that runs less than 30 times, it will not optimize even after 1000 runs. Also sometimes we saw inconsistencies even in this. It seems that maybe Chrome is using shared counters across functions or is triggering some other limit. However our counters do affect the optimization. For example a high loop counter will cause it to always optimize anything with any loop. And if there is no loop, our code will optimize after whatever number the user picked runs.

IV. IMPLEMENTATION

The implementation of these optimizations should have been fairly simple. However V8 is a heavily regimented and organized code-base and we will explain the difficulties with this in the next section. In V8 there is a structure called the *Isolate*. The *Isolate* is a representation for a thread in V8 or when used by Chrome, a 'tab' in the Chrome browser. It has it's own heap and garbage collection. In our code we use the *Isolate* as similar to a global state.

Our implementation first searches the first three characters of all multi-line comments looking for the signature 'LEZ'. If this is found, it calls our code to begin parsing. It then adds this code into the *Isolate* (we added implementation to access the isolate at this point). We store a map in the *Isolate* with the key being function name and the value being a pointer to a malloc-ed int *. The first int in the array is a bit map where each bit corresponds to whether an optimization is 'on' or 'off'. The other spots in the array are for extra information, for example the user input in the Edge Weight Profile optimization. The locations of these values are hard-coded as defines to make them easy to access.

The next part was fairly tricky as we wanted to avoid having to do string based map lookups every time we wanted to find out if a function needed to be optimized. So we solved this problem by attaching the optimization data to a per-function data structure. There were several options: Code object, *JSFunction* object, and *SharedFunctionInfo* objects. We decided to add it to *SharedFunctionInfo* because this object had a lot of information stored with it and was available from the *JSFunction* object. Also Code and *JSFunction* objects seem to be recreated and deleted often. We copied the int array into a structure that was used inside V8 to avoid lots of issues. This was done on creation of the *SharedFunctionInfo* and the information was pulled from the isolate with a map lookup. Note that we do have to do a hash map lookup once for every function even if it is not optimized.

Once we had the data inside the *SharedFunctionInfo* object we could use it to actually make the optimizations. For the *Optimize Always* optimization we made it do the optimization immediately by modifying a check in V8 that optimizes if the V8 flag *always_opt* is on. To implement *Never Optimize*, we tricked V8 into thinking that it shouldn't optimize the function, by setting a reason in

the *SharedFunctionInfo* to deoptimize the function. For the third optimization, Edge Weight Profile, we modified code that writes assembly code to set the previously discussed weight value of the counters. We instead, made it look for the values from the flag-object we added to *SharedFunctionInfo*.

Another thing we implemented was to add a flag to turn our optimizations on and off. If the flag indicates not to use our optimizations, no parsing is done and a lookup into the map does not happen and no checks to the object in *SharedFunctionInfo*. This should allow us to test our implementation's impact and test the optimizations compared to none.

V. DIFFICULTIES

We ran into many difficulties with V8. V8 has a very complex and regimented codebase with very few comments or documentation regarding more than a general overview. Here are a few of the things we had to struggle with in V8.

For one, *SharedFunctionInfo*, *JSFunction*, and *Code* are very regimented data structures. Their data is accessed by static offsets and their total size is used. Also there are requirements on their data that it be a heap object and it seems to create some sort of snapshot on compile and has a lot requirements regarding that. Also since V8 runs it's own code during compile, it would cause segmentation faults with no usable output.

Another difficulty with V8 is that it takes forever to compile. It took 30 minutes to compile and link a header file and on our computer and 15 minutes to even compile and link a change to a cc file.

We also struggled with the V8 garbage collector as they have special handle objects to handle garbage collected objects. This is to be able to change pointers when sweeping the objects. However these could not be put into a map and caused lots of issues which is why we had to use a simple malloced int[] and copy it into a V8 heap object when copying into the *SharedFunctionInfo*.

We ran into a lot of other issues with V8 including the Printf (that they use for logging) not always actually printing out to stdout and dealing with Chrome caching web pages on load.

Another issue we had with Chrome was that it has so many pieces to the optimization that it was hard to figure out exactly what it was doing. Usually it seemed to follow simple rules but sometimes if you did things like add more functions it would no longer follow them. This made it hard to test our optimizations and the effect they had on Chrome.

VI. ANALYSIS AND RESULTS

First, we will explain the timing method used here. We have placed separate timers to measure the full-code

compile time and optimized compile for every function. We have also placed timers to track the complete execution time. This should allow us to evaluate and prove that each of our optimizations works. We also observed our overhead by timing how long it took to parse the extra code and the time it took to do the map lookups. This was very little overhead however we will include it in our results. All numbers are based off an average and standard deviation of 6 runs. The benchmarks were all run on the same laptop with 8GB memory and 2.7 GHz i5 quad core.

A. Opt 1: Optimize Immediately

We verified that this optimization worked by picking a simple javascript that would not normally be optimized in V8. The example we used is a simple function looped 10 times. V8 logged (when run with the flag `-trace_opt_verbose`) that it had optimized the function and our supplemental logging backed up our results. However we noticed upon further analysis that it was deoptimizing the code immediately when it first ran it. On further inspection the reason for this was that the function did not have enough information. This makes sense because before a function is run you know nothing about the types. However we are not sure why Chrome chose to optimize successfully before it deoptimized instead of just failing to optimize. To test our hypothesis, we attempted to optimize a function that just did computations on constants. V8 then optimized the code and never deoptimized it. This makes sense because it knows all the information regarding the function. On further analysis, it will deoptimize if any variable is used on the right hand side of an assignment. We were also able to print out the reason for deoptimization in the previous code and it was most often "Insufficient type feedback for LHS of binary operation", though we also saw other issues in more complex code.

Because the limitations of this optimization are not very interesting, we did not do further analysis. However, one can fairly easily implement Optimize immediately by using the edge weight profiling optimization, as explained below.

B. Opt 2: Never Optimize

This optimization was the easiest to implement and test. V8 prints out that it wanted to optimize a function but couldn't and prints out the reason that we gave it and we can see that it never optimizes the function in the text output. However we also ran a benchmark to prove that our deoptimization works. We created a function that had a for-loop that ran 100000 times inside it. Then we ran that function 2000 times. We compared a run where we told the function not to optimize and

run that ignored our extensions and optimized.

| | Unmodified | Modified |
|------------------|------------------|-------------------|
| Execution | 1525.2 ± 82.3 ms | 2731.7 ± 195.9 ms |
| Compile | 87.65 ± 8.7ms | 101.5 ± 15.0 ms |
| Overhead | 0 | .428 ± .425 ms |

The original run was a lot faster than our run because that function was optimized. This shows that our deoptimization function works.

Next we present a case where our optimization actually causes faster execution. To test this, we create 9 functions that did simple operations including a small while loop of 100. We then ran each of these functions 135 times. This would cause the functions to optimize in the original case and not optimize in our case. This should save us on optimization time.

| | Unmodified | Modified |
|---------------------|-------------------|----------------------|
| Execution | 1744.1 ± 246.9 ms | 1600.7 ± 207.4 ms |
| Opt Compile | 85.8 ± 17.7ms | 58.72 ± 16.8 ms |
| Full Compile | 83.55 ± 12.8 | 93.03 ± 14.14ms |
| Overhead | 0 | .837 ± .810 ms |
| Sum | 1913.4 ± 231.7 ms | 1753.317 ± 231.69 ms |

These results are a little confusing. We can see that we gain some time on the fact that we don't optimize. However we seem to be faster for execution as well. One pattern that we noticed was that execution seemed to take a little longer when an optimization was done. We are not sure why this is. It is possible it is the extra cost of setting up the optimization. However one may notice that it is only around 1 standard deviation higher, so it is not that much. Even with just the improvement in compile time, this optimization would be successful.

C. Opt 3: Edge Weight Profiling

We tested that this optimization works by running many different kinds of codes and verifying that they optimized when normally they wouldn't.

A good use of this optimization is if your function runs just enough times to be optimized. In this case default V8 would optimize but then quickly end and not get a good result. Our test case had 20 similar functions. Each function did a few simple calculations and a while loop that ran to 30. Each function was run 550 times. We gave the first input as 10 (to optimize faster) and the second input as 5 (to optimize on lower number of while loops). This should cause our function to optimize very quickly. The results are shown below.

| | Unmodified | Modified |
|------------------|------------------|-------------------|
| Execution | 1217.4 ± 70.5 ms | 1126.9 ± 44.8 ms |
| Compile | 99.4 ± 9.3 ms | 101.8 ± 8.1 ms |
| Overhead | 0 | 0.583 ± 0.54 ms |
| Sum | 1316.8 ± 78.2 ms | 1229.24 ± 48.8 ms |

As you can see, both default and ours optimize. However we have less execution time because our optimization

was done earlier.

VII. WHAT WE LEARNED

One of the big things that was interesting to see was how V8 worked compared to the compilers we have talked about in class. It was interesting to see things we had learned in class in the V8 engine, and also how it was different. For example how their garbage collector uses Handles as was mentioned in the slides. It was also interesting to see how they actually implemented edge profiling and how they handled deoptimization and optimization and the different heuristics that were used. It is likely that these heuristics went through a lot of testing to arrive where they are now.

VIII. LIMITATIONS

There are a few limitations with our implementation. One is that we currently just place our function names into the *Isolate*. This means that these function names will be loaded on future pages that are loaded in that

tab. This is a performance issue (if too many people use it it will make the stored data larger) and also a security issue as it allows a page to slightly influence other pages that are loaded.

Another issue is that Chrome's optimizations are very complex and we were unable to test infinite test cases. We know our optimizations work for simple cases but they might not be as effective for larger files.

A note here, that all our tests were run on initial page loads. However V8 can cache code between runs so our optimizations may not help as much if a user were to refresh a page.

IX. CONCLUSION

All in all, our optimizations show that some performance gains can be made from allowing user and developers to add their own hints to the compiler. With better knowledge of V8, we can improve the performance even more, by tapping aggressive flags. This would be a future course of action. However, it seems V8 does do a really good job of guessing the optimizations itself.
